

Junze Zhu

Ph.D. Candidate | Artificial Intelligence, Nanjing University

aminataconde675@gmail.com | huicir.github.io | Google Scholar | GitHub



Research interests

Language agents, multi-agent systems, speculative and asynchronous execution, harness-aligned policy learning, knowledge graphs, and recommendation.

Education

2025–2029 (expected) **Ph.D. Candidate in Artificial Intelligence**

Nanjing University
Focus: Agent Systems and Swarm Intelligence

2022–2025 **Master's in Data Science**

Wuhan University
Focus: Knowledge-Enhanced Intelligent Systems

2018–2025 **Bachelor's in Electronic Information**

Wuhan University
Focus: Multimodal Fusion Perception

Selected publications

2026 Recognize Your Orchestrator: An Entropy Dynamics Perspective for LLM Multi-Agent Systems
Junze Zhu, Weihao Chen, Xuanwang Zhang, Zhen Wu, Xinyu Dai. *ICML 2026*, PMLR 306:166844–166869. Paper.

Mean-field entropy dynamics for understanding instability in centralized language-agent orchestration.

2025 Knowledge Graph Preference Contrastive Learning for Recommendation
Junze Zhu, Zhongyi Hu, Fan Zhang. *Learning on Graphs Conference*, PMLR 269:18:1–18:15. Paper / Code.
Preference-aware contrastive learning and rationale attention for knowledge-aware recommendation.

2026 WebProgramBench: Can Coding Agents Reconstruct Live Applications from Scratch via Black-Box Interaction?

Xuanwang Zhang, Yuteng Han, Xinyu Liu, **Junze Zhu**, et al. *ICLR 2027 submission*, under review; fourth author. OpenReview.

A benchmark for executable reconstruction of live applications through black-box interaction.

Manuscripts and research projects

- **Anchoring and Adjustment: A Cognitive Refinement Framework for Trustworthy Retrieval Augmentation.** Working manuscript, 2026.
- **Perseus: Non-Blocking Parallel Exploration with a Speculative Swarm.** Under review. An asynchronous agent swarm architecture that expands inference-time exploration while an authoritative Actor retains task control. Project site.
- **Harness Policy Distillation: Aligning Agent Policies with Runtime Semantics.** Under review. A policy learning framework that distills verified actions and runtime control-document views into a harness-compatible agent policy.

Selected systems work

- **Open Crowd** – skill and CLI management for complex task orchestration with agent swarms.
- **Token-Never-Sleeps** – long-horizon agent swarm task management and external-loop execution; published as an npm package.

Links

Academic homepage | GitHub | Google Scholar | WebProgramBench on OpenReview